

KALPESH KRISHNA

kalpeshk2011@gmail.com ♦ [LinkedIn](#) ♦ [Github](#) ♦ [Twitter](#) ♦ [Google Scholar](#) ♦ [Semantic Scholar](#)

WEBPAGE: <https://martiansideofthemoon.github.io>

EMPLOYMENT

Google DeepMind (GenAI, Gemini)

Staff Research Scientist

Senior Research Scientist

Focus Areas: LLM instruction following, RL, evals

Oct '25 - Present

Nov'24 - Oct '25

New York, NY

Google (Gemini App)

Research Scientist

Focus Areas: LLM instruction following, RL, evals

August '23 - Nov '24

New York, NY

EDUCATION

University of Massachusetts, Amherst

MS/PhD in Computer Science

(advised by [Prof. Mohit Iyyer](#))

Major GPA: 4.0/4.0

Sept '18 - Sept '23

Indian Institute of Technology, Bombay

B.Tech in Electrical Engineering, Minor in Computer Science

(advised by [Prof. Preethi Jyothi](#))

Major GPA: 9.74/10 (2nd among 66 students)

Minor GPA: 10/10

July '14 - July '18

INTERNSHIPS

Allen Institute for Artificial Intelligence (AI2)

Research Intern under [Kyle Lo](#), [Arman Cohan](#), [Pradeep Dasigi](#)

June '22 - August '22

Seattle, WA

Google DeepMind (formerly Google Research)

Student Researcher under [John Wieting](#)

October '21 - May '22

Amherst, MA (remote)

Google DeepMind (formerly Google Research)

Research Intern under [Partha Talukdar](#) and [Bidisha Samanta](#)

June '21 - September '21

Bangalore, India (remote)

Google DeepMind (formerly Google Brain)

Research Intern under [Aurko Roy](#)

May '20 - August '20

Mountain View, CA (remote)

Google DeepMind (formerly Google Research)

Research Intern under [Gaurav Singh Tomar](#) and [Ankur Parikh](#)

May '19 - September '19

New York, NY

Toyota Technological Institute at Chicago

Research Intern under [Karen Livescu](#), [Liang Lu](#) and [Kevin Gimpel](#)

May '17 - July '17

Chicago, IL

Mozilla Foundation

Google Summer of Code Intern under [Armen Zambrano](#)

May '16 - August '16

Mumbai, India / London, UK

AWARDS & SCHOLARSHIPS

- Google PhD Fellowship 2021-23 ([list of recipients](#))
- Paper Awards:
Distinguished Paper Award, ACM CCS 2023 (5 out of 800 submissions)
Outstanding Paper Award - EACL 2023 (12 out of 1600 submissions)
- Sharad Maloo Memorial Gold Medal at IIT Bombay
(For outstanding academic and extra-curricular achievements)
- Academic service awards:
Outstanding Area Chair at EMNLP 2024 (top 6% area chairs)
Outstanding Reviewer at NeurIPS 2025; ICLR 2021-22; ACL 2020, 2022 (Top 7-12% reviewers)
- PhD Candidacy with Distinction at UMass Amherst
- Cargill Global Leaders Scholarship 2016-18
(Awarded by the [IIE](#) and [Cargill](#) for academic, extra-curricular and leadership achievements)

INVITED TALKS

- Georgia Tech, April 2024
(Towards robust long-form text generation evaluation)
- University of Texas at Dallas, November 2023
(Towards robust long-form text generation evaluation)
- University of Pittsburgh, November 2023
(Towards robust long-form text generation evaluation)
(Paraphrasing evades detection of AI-generated text, but retrieval is an effective defense)
- NLP with Friends, September 2023
(Paraphrasing evades detection of AI-generated text, but retrieval is an effective defense)
- IBM Research, July 2023
(Towards robust long-form text generation evaluation)
- University of Toronto, July 2023
(Paraphrasing evades detection of AI-generated text, but retrieval is an effective defense)
- Allen Institute for Artificial Intelligence, May 2023
Google Deepmind, May 2023
NVIDIA Applied Deep Learning Research, April 2023
Apple Machine Learning Research, March 2023
(Towards robust long-form text generation systems)
- University of Washington, September 2022
(RankGen: Improving Text Generation with Large Ranking Models)
- University of Southern California, July 2021
University of Texas at Austin, June 2021
(Progress in text generation & perils of its evaluation)

- Google Research, May 2021
(*Hurdles to progress in long-form question answering*)
- IBM Research, April 2020
(*On model extraction attacks on BERT-based APIs*)
- AllenNLP Summit at AI2, August 2019
(*Lightning talks on using AllenNLP for [education](#)*)
- UMass Data Science Research Symposium, April 2019
(*On hierarchical question-answer generation*)

ACADEMIC SERVICE

- **Senior Area Chair** for ACL 2025 (Track: Language Modeling, [full committee](#))
- **Area Chair** for 50+ papers in EMNLP 2024-26, ACL 2024, NAACL 2024-25, EACL 2024-26 and workshops
- **Program Committee / Reviewer** for 100+ total papers in EMNLP '19-'23; ACL '19-'23; ICLR '21-'25; NAACL '19-'22; NeurIPS '21-'25; ICML '23-'24, EACL '23; IJCV '19; NEJLT '23; COLM '24-'25 and several workshops in NeurIPS, ACL, NAACL, EMNLP
Outstanding / Top Reviewer at NeurIPS 2025; ICLR 2021-22; ACL 2020, 2022; (Top 7-12% reviewers)
- **Student Volunteer** at ACL '19 (awarded [ACL 2019 Student Scholarship](#)), AKBC '19, ICLR '20

TEACHING / MENTORING / OTHER SERVICE

- **Interns:**
[Anthony Liang](#) (USC) at Google DeepMind (Summer 2025)
[Satyapriya Krishna](#) (Harvard University) at Google Gemini (Summer 2024)
- **Mentored** five UMass undergraduates on their honors thesis (2019-2021) and four student teams for the industrial mentorship projects (COMPSCI 696) in Spring 2019-2022
- **Co-organized** a weekly talk series “Machine Learning and Friends Lunch” at UMass (2019-22)
- **Graduate Teaching Assistant** at UMass Amherst for *Deep Learning for NLP* (Spring 2019) and *Advanced Natural Language Processing* (Fall 2020)
- **Institute Student Mentor** at IIT Bombay (2017-18)
- **Manager, Web and Coding Club** at IIT Bombay (2016-17)
Awarded the Institute Organizational Color 2016-17 at IIT Bombay for achievements while leading the Web and Coding Club
- **Teaching Assistant** at IIT Bombay in *Computer Programming* (2016) and *Linear Algebra* (2017)
- **Web Coordinator** for Mood Indigo 2016 at IIT Bombay

As per [Google Scholar](#), these papers have been **cited 14400 times** (as of October 2025).
NeurIPS, ICLR, EMNLP, *ACL are top-tier peer-reviewed conferences in ML/NLP with an acceptance rate of 20-30%.

Selected + Recent Papers:

* denotes equal contribution / co-lead authors.

1. [Gemini 2.5: Pushing the Frontier with Advanced Reasoning, Multimodality, Long Context, and Next Generation Agentic Capabilities](#)
Gemini Team, Google DeepMind (*Kalpesh Krishna* is co-author and core contributor)
arXiv 2025
2. [Fact, Fetch, and Reason: A Unified Evaluation of Retrieval-Augmented Generation](#)
Satyapriya Krishna, *Kalpesh Krishna**, Anhad Mohananey*, Steven Schwarcz, Adam Stambler, Shyam Upadhyay, Manaal Faruqui
arXiv 2024
3. [Foundational Autoraters: Taming Large Language Models for Better Automatic Evaluation](#)
Tu Vu*, *Kalpesh Krishna**, Salaheddin Alzubi, Chris Tar, Manaal Faruqui, Yun-Hsuan Sung
EMNLP 2024
4. [Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context](#)
Gemini Team, Google (*Kalpesh Krishna* listed as a Core Contributor)
arXiv 2024
5. [Gemini: A family of highly capable multimodal models](#)
Gemini Team, Google (*Kalpesh Krishna* listed as a Core Contributor)
arXiv 2024
6. [FActScore: Fine-grained Atomic Evaluation of Factual Precision in Long Form Text Generation](#)
Sewon Min*, *Kalpesh Krishna**, Xinxi Lyu, Mike Lewis, Wen-tau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, Hannaneh Hajishirzi
EMNLP 2023
7. [Paraphrasing evades detectors of AI-generated text, but retrieval is an effective defense](#)
Kalpesh Krishna, Yixiao Song, Marzena Karpinska, John Wieting, Mohit Iyyer
NeurIPS 2023
8. [LongEval: Guidelines for Human Evaluation of Faithfulness in Long-form Summarization](#)
Kalpesh Krishna, Erin Bransom, B. Kuehl, Mohit Iyyer, Pradeep Dasigi, Arman Cohan, Kyle Lo
EACL 2023 (**Outstanding Paper Award**)
9. [Hurdles to Progress in Long-form Question Answering](#)
Kalpesh Krishna, Aurko Roy, Mohit Iyyer
NAACL 2021
10. [Reformulating Unsupervised Style Transfer as Paraphrase Generation](#)
Kalpesh Krishna, John Wieting, Mohit Iyyer
EMNLP 2020

11. [Thieves on Sesame Street! Model Extraction of BERT-based APIs](#)
Kalpesh Krishna, Gaurav S. Tomar, Ankur P. Parikh, Nicolas Papernot, Mohit Iyyer
ICLR 2020

Other Papers:

12. [PostMark: A Robust Blackbox Watermark for Large Language Models](#)
Yapei Chang, *Kalpesh Krishna*, Amir Houmansadr, John Wieting, Mohit Iyyer
EMNLP 2024
13. [GEE! Grammar Error Explanation with Large Language Models](#)
Yixiao Song, *Kalpesh Krishna*, Rajesh Bhatt, Kevin Gimpel, Mohit Iyyer
NAACL 2024 (Findings)
14. [On the Risks of Stealing the Decoding Algorithms of Language Models](#)
Ali Naseh, *Kalpesh Krishna*, Mohit Iyyer, Amir Houmansadr
CCS 2023 (Distinguished Paper Award)
15. [ezCoref: Towards Unifying Annotation Guidelines for Coreference Resolution](#)
Ankita Gupta, Marzena Karpinska, Wenlong Zhao, *Kalpesh Krishna*, Jack Merullo, Luke Yeh, Mohit Iyyer, Brendan O'Connor
EACL 2023 (Findings)
16. [NL-Augmenter: A Framework for Task-Sensitive Natural Language Augmentation](#)
Kaustubh D. Dhole and others
NEJLT 2023
17. [RankGen: Improving Text Generation using Large Ranking Models](#)
Kalpesh Krishna, Yapei Chang, John Wieting, Mohit Iyyer
EMNLP 2022
18. [SLING: Sino Linguistic Evaluation of Large Language Models](#)
Yixiao Song, *Kalpesh Krishna*, Rajesh Bhatt, Mohit Iyyer
EMNLP 2022
19. [Document-Level Literary Machine Translation with Parallel Paragraphs from World Literature](#)
Katherine Thai*, Marzena Karpinska*, *Kalpesh Krishna*, Bill Ray, Moira Inghilleri, John Wieting and Mohit Iyyer
EMNLP 2022
20. [Few-shot Controllable Style Transfer for Low-Resource Multilingual Settings](#)
Kalpesh Krishna, Xavier Garcia, Deepak Nathani, Bidisha Samanta, Partha Talukdar
ACL 2022
21. [RELiC: Retrieving Evidence for Literary Claims](#)
Katherine Thai, Yapei Chang, *Kalpesh Krishna*, Mohit Iyyer
ACL 2022
22. [User and Technical Perspectives of Controllable Code Generation](#)
Stephanie Houde, Vignesh Radhakrishna, Praneeth Reddy, Juie Darwade, Haoran Hu, *Kalpesh Krishna*, Mayank Agarwal, Kartik Talamadupula, Justin D. Weisz
NeurIPS HCAI workshop 2022 (short paper)

23. [Do Long-Range Language Models Actually Use Long-Range Context?](#)
Simeng Sun, *Kalpesh Krishna*, Andrew Mattarella-Micke and Mohit Iyyer
EMNLP 2021
24. [Weakly-Supervised Open-Retrieval Conversational Question Answering](#)
Chen Qu, Liu Yang, Cen Chen, W. Bruce Croft, *Kalpesh Krishna* and Mohit Iyyer
ECIR 2021
25. [Long Document Summarization in a Low Resource Setting using Pretrained Language Models](#)
Ahsaas Bajaj, Pavitra Dangati, *Kalpesh Krishna*, Pradhiksha Kumar, Rheeya Uppaal, Bradford Windsor, Eliot Brenner, Dominic Dotterer, Rajarshi Das, Andrew McCallum
ACL Student Research Workshop 2021
26. [An Analysis of Frame-skipping in Reinforcement Learning](#)
Shivaram Kalyanakrishnan, Siddharth Aravindan, Vishwajeet Bagdawat, Varun Bhatt, Harshith Goka, Archit Gupta, *Kalpesh Krishna*, Vihari Piratla
arXiv 2021
27. [SunPy: A Python package for Solar Physics](#)
Stuart J. Mumford and *others*
JOSS 2020
28. [Generating Question-Answer Hierarchies](#)
Kalpesh Krishna, Mohit Iyyer
ACL 2019
29. [Syntactically Supervised Transformers for Faster Neural Machine Translation](#)
Nader Akoury, *Kalpesh Krishna*, Mohit Iyyer
ACL 2019
30. [Trick or TReAT: Thematic Reinforcement for Artistic Typography](#)
Purva Tendulkar, *Kalpesh Krishna*, Ramprasaath R. Selvaraju, Devi Parikh
ICCC 2019 (**Best Presentation Award**)
31. [Revisiting the Importance of Encoding Logic Rules in Sentiment Classification](#)
Kalpesh Krishna, Preethi Jyothi, Mohit Iyyer
EMNLP 2018 (*short paper*)
32. [Hierarchical Multitask Learning for CTC-based Speech Recognition](#)
Kalpesh Krishna, Shubham Toshniwal, Karen Livescu
arXiv 2018
33. [A Study of All-Convolutional Encoders for Connectionist Temporal Classification](#)
Kalpesh Krishna, Liang Lu, Kevin Gimpel, Karen Livescu
ICASSP 2018 (*Awarded SPS Travel Grant*)

Other Projects: https://martiansideofthemoon.github.io/other_projects/

OTHER ACHIEVEMENTS

- CICS Graduate Fellowship 2018, Victor Lesser Graduate Scholarship 2019 at UMass Amherst
- Selected for [JSALT '17](#), organized by JHU's [Center for Language and Speech Processing](#)²
- Institute Academic Prize 2015-16 at IIT Bombay
- Top 10 at the Astronomy Olympiad's Indian Selection Camp for IOAA '14, (*20000 candidates*)
- All India Rank 93 in [JEE Advanced](#) '14 (*out of 126,000 candidates*)
All India Rank 34 in [JEE Mains](#) '14 (*out of 1,400,000 candidates*)
All India Top 100 in [CBSE XII](#) '14 (*out of 1,000,000 candidates*)
All India Rank 2 in [ICSE](#) '12 (*out of 132,000 candidates*)

Awarded by Pune city institutions for being the city topper in each of these exams. Featured in the Pune city editions of several newspapers including The Times of India and The Indian Express.

- Passed the [National Standard Exam in Physics](#) '14 and [National Standard Exam in Chemistry](#) '14, the first rounds in the Indian selection procedure for the International Physics Olympiad and the International Chemistry Olympiad respectively
- Selected for the [Kishore Vaigyanik Protsahan Yojana](#) (KVPY) Award '14 (1,000/20,000 applicants)

EXTRA-CURRICULAR

- Board Member of the Hampshire Bird Club in Massachusetts (2022-24)
 - **Blogging** - Active [blogger](#) on MS/PhD applications, biking, bird photography and computer science. Led development of a [beginner's programming wiki](#) with the Web and Coding Club, IITB
 - **StackOverflow** - Active [contributor](#) in 2016-17, top 6% overall
 - **Open Source** - Actively [contributed](#) to Mozilla in 2015-16 including [QoC 2015](#), [GSoC 2016](#)
 - **Karate** - Black Belt (1st Dan) trained in Kyokushin Kai for seven years, district level winner
 - **Abacus & Mental Arithmetic** - [Aloha](#) grandmaster, national and state level winner
- Hobbies** — [birdwatching](#), macro & bird photography, hiking, biking, reading, [star gazing](#)

²Couldn't attend due to clashing college schedule